



NVIDIA's Brev Platform

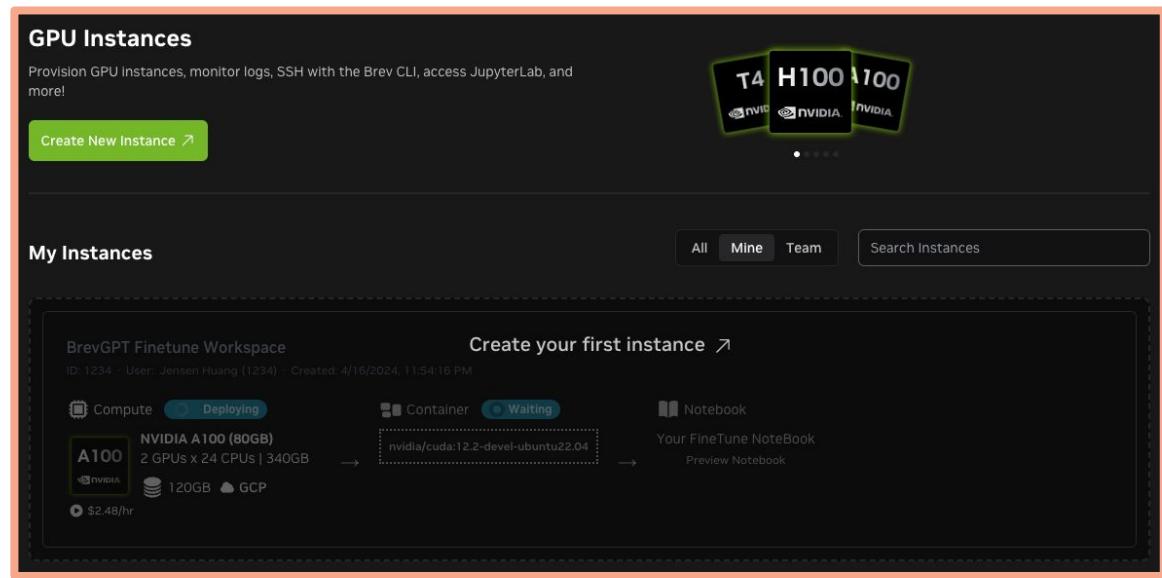
Develop with the latest NVIDIA GPUs and Software

Introduction to Brev

Your starting point for AI development and experimentation

NVIDIA's Brev platform simplifies AI development by aggregating multiple cloud providers to deliver the right NVIDIA GPUs and environments. Developers can quickly launch and customize GPU-backed setups without complex infrastructure management.

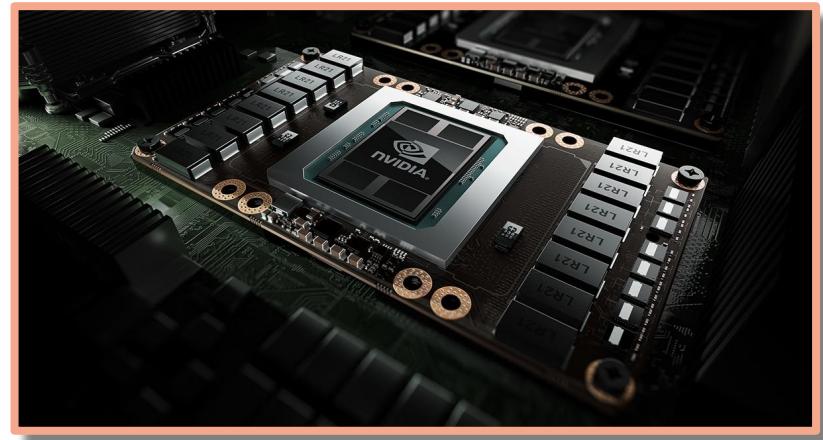
- **Instant GPU access:** On-demand, high-performance NVIDIA GPUs provisioned instantly through our Partners. Virtual Machines preconfigured with *Python*, *CUDA*, *Docker*, and *Jupyter Notebooks*
- **Launchables:** One-click, customizable GPU environments for faster onboarding and team consistency across projects
- **CLI & IDE Support:** Access via Brev CLI and integrate with Cursor, Windsurf, VS Code, and tmux



Access to Leading Cloud Providers and NVIDIA GPUs

Constantly Growing...

NVIDIA's Brev connects you to a worldwide NVIDIA Cloud Partner network, delivering cutting-edge GPUs on demand for AI/ML acceleration.



Pricing & Packaging

Self-service, pay-as-you-go GPU access for maximum flexibility

- **Flexible Pay-as-You-Go:** Purchase credits to access on-demand GPU compute, ensuring cost-efficiency with no long-term commitments.
- **Transparent Hourly Rates:** Clear pricing per GPU instance type, aggregated from our Partners

Select your Compute

VM Mode w/ Jupyter | Configure | VRAM (GPU Memory) Any Size | Cloud Brev Cloud - any | Filters

B200 H200 H100 A100 L40S A10 A10G

1x NVIDIA H200

1x NVIDIA H200 150GiB VRAM 200GiB RAM x 16 CPUs 1TiB SSD SXM NEBIUS \$4.20/hr

The screenshot shows a user interface for selecting a GPU instance. At the top, there's a header 'Select your Compute' with several filter options: 'VM Mode w/ Jupyter' (selected), 'Configure', 'VRAM (GPU Memory) Any Size', 'Cloud Brev Cloud - any', and 'Filters'. Below this, there's a row of GPU icons: B200, H200, H100, A100, L40S, A10, A10G. The H200 icon is highlighted with a green glow, indicating it's selected. Below the icons, it says '1x NVIDIA H200'. At the bottom, there's a detailed breakdown of the instance: '1x NVIDIA H200 150GiB VRAM 200GiB RAM x 16 CPUs 1TiB SSD SXM NEBIUS \$4.20/hr'. The NVIDIA logo is visible in the bottom right corner.

NVIDIA *Launchables*

“

A one-click experience — no setup, just code running on NVIDIA’s Accelerated Computing stack. It’s a frictionless way to show, share, and scale innovation.

”

Launchables

1-click deployments for pre-configured software and hardware combinations

The screenshot shows a detailed view of a pre-configured software and hardware combination. At the top, the title is "single-cell-analysis-GTC-Paris" with a price of "\$1.74/hr". Below this, there's a "Compute" section showing an "L40S" node with "NVIDIA L40S (48GB)" memory, "1 GPU" and "8 CPUs", and "147GB" storage. A green arrow labeled "Price" points to the price information. Another green arrow labeled "Predefined Compute" points to the "Compute" section. To the right, a large green arrow labeled "Software is pre-configured" points to the "Container" section, which displays a GitHub repository URL: "github.com.../docker-compose-nb-2504.yaml". Further down, there's a "Files" section showing a GitHub repository: "GitHub Repo: single-cell-analysis-blueprint" with a "View Source" link. A green arrow labeled "Content Preview" points to this section. Below these, there's a "Repository Info" section for "single-cell-analysis-blueprint" by NVIDIA-AI-Blueprints, showing "17 stars", "10 forks", and "17 watchers". In the "README.md" section, there are two purple buttons: "RAPIDS" and "scverse". A green arrow labeled "Content Preview" also points to this section. The bottom part of the screenshot contains a detailed description of the repository's purpose: "This repository houses tutorial notebooks to run GPU-accelerated single-cell analysis workflows using RAPIDS-singlecell, a GPU accelerated library developed by scverse®. The goal is of this repository is to help users try out and explore different capabilities of RAPIDS-singlecell on datasets ranging from 250 thousand to 11 million cells. To make this as easy as possible, we set up two different GPU environments on Brev that are designed to get you working with GPU-". The NVIDIA logo is in the bottom right corner.

single-cell-analysis-GTC-Paris
\$1.74/hr

Compute

L40S NVIDIA L40S (48GB)
1 GPU x 8 CPUs | 147GB

Container

github.com.../docker-compose-nb-2504.yaml

Files

GitHub Repo: single-cell-analysis-blueprint
View Source

Exposed Ports

Tunnels: jupyter:8888 |
disk1:8787 |
disk2:8786

Content Preview:

This repository will be present in the instance after deployment

Repository Info

single-cell-analysis-blueprint by NVIDIA-AI-Blueprints

17 stars 10 forks 17 watchers

README.md

RAPIDS

scverse

Single-Cell Analysis Blueprint

This repository houses tutorial notebooks to run GPU-accelerated single-cell analysis workflows using RAPIDS-singlecell, a GPU accelerated library developed by scverse®. The goal is of this repository is to help users try out and explore different capabilities of RAPIDS-singlecell on datasets ranging from 250 thousand to 11 million cells. To make this as easy as possible, we set up two different GPU environments on Brev that are designed to get you working with GPU-

Sign in

Deploy Launchable

NVIDIA



Walkthrough+Demo

- *Create Account*
- *Create Org*
- *Redeem Coupon (if there is one)*
- *Invite Users*
- *2 ways to start:*
 - *1-Click Deploy with Launchables*
 - *Create an instance from scratch*

Happy Coding!

- **Brev Docs & Guides:** <https://docs.nvidia.com/brev/latest/about-brev.html>
- **GPU Catalog:** build.nvidia.com/gpu
- **Platform or Technical Support:** brev-support@nvidia.com